

DS5559-Final_proj-NHAMCS-Report_2

June 28, 2020

1 Final Project: Admission Prediction from NHAMCS

1.1 Progress report: Initial model evaluation

1.1.1 DS5559: Big Data Analysis

1.1.2 Thomas Hartka(trh6u), Alicia Doan(ad2ew), Michael Langmayr(ml8vp)

Created: 6/28/20

In this report we demonstrate a logistic regression model. We use patient characteristics as input and create a model to predict hospital admission. We transform our variables and create a vector of features. We then loop through values of α for standard, ridge, and lasso methods. All the models use weights for the outcomes since there is class imbalance. The best hyperparameters were selected based on AUC.

The following is a summary of the best LR model:

- Type of model:
 - **Logistic Regression**
- Best hyperparameters used:
 - **Method: Ridge regression**
 - **Regularization parameter: 1.0**
- Size of the saved model:
 - **Disk usage: 52k**
- Performance metrics:
 - **Accuracy: 0.713**
 - **Area under ROC curve (AUROC): 0.755**
 - **F1 score: 0.183**
 - **Confusion matrix:**
tn: 1473 fn: 45
fp: 587 tp: 99

1.2 Configuration

```
[1]: # set data directory
data_dir = "../data"
results_dir = "../results"
```

1.3 Import libraries and set up Spark

```
[2]: # import python libraries
import os
import pandas as pd
import numpy as np
from functools import reduce
```

```
[3]: # set up pyspark
from pyspark.sql import *
from pyspark.sql import SparkSession
from pyspark.sql.functions import *
from pyspark.sql.types import IntegerType
```

```
[4]: from pyspark.ml import Pipeline
from pyspark.ml.feature import *
from pyspark.ml.classification import LogisticRegression
from pyspark.ml.evaluation import BinaryClassificationEvaluator
```

```
[5]: spark = SparkSession.builder.getOrCreate()
```

1.4 Read in data

```
[6]: NHAMCS = spark.read.parquet(data_dir + "/NHAMCS_processed.2007-2017")
```

1.5 Transform data

```
[7]: # perform string indexing to prepare for OHE for residence variable
rsi = StringIndexer(inputCol="RESIDNCE", outputCol="RESINDEX")
simodel = rsi.fit(NHAMCS)
NHAMCS = simodel.transform(NHAMCS)
```

```
[8]: # perform OHE on residence variable
rohe = OneHotEncoder(inputCol='RESINDEX', outputCol='RESONE')
NHAMCS = rohe.transform(NHAMCS)
```

```
[9]: # assemble vector
va = VectorAssembler(inputCols=["AGEYEAR", "RESONE", 'SEXMALE', 'ARRTIMEMIN', 'YEAR', 'PULSE', 'TEMPF',
    'RESPR', 'BPSYS', 'BPDIAS', 'POPCT', 'PAINSCALE', 'ALZHD', 'ASTHMA', 'CAD', 'CANCER',
    'CEBVD', 'CHF', 'CKD', 'COPD', 'DEPRN', 'DIABTYP0', 'DIABTYP1', 'DIABTYP2', 'EDHIV',
    'ESRD', 'ETOHAB', 'HPE', 'HTN', 'HYPLIPID', 'OBESITY', 'OSA', 'OSTPRISIS', 'SUBSTAB',
    'NOCHRON', 'TOTCHRON', 'INJURY', 'INJURY72'],
    outputCol="features")

NHAMCS = va.setHandleInvalid("skip").transform(NHAMCS)
```

1.6 Train and test model

```
[10]: # split into training and testing set
training, testing = NHAMCS.randomSplit([0.8, 0.2], 42)
```

```
[11]: # handle class imbalance

# calculate balance ratio
balRatio = training.select("ADM_OUTCOME").where('ADM_OUTCOME == 0').count() /
    training.count()

# add weights
training = training.withColumn("classWeights", when(training.ADM_OUTCOME ==
    1, balRatio).otherwise(1-balRatio))
```

```
[12]: # function for logistic regression
def lr_nhamcs (training_set, testing_set, reg_param=0, method="Standard"):
    if method=="Standard":
        lr = LogisticRegression(featuresCol="features", labelCol="ADM_OUTCOME",
            weightCol="classWeights", \
                maxIter=10, regParam=0, elasticNetParam=0)
    elif method=="Ridge":
        lr = LogisticRegression(featuresCol="features", labelCol="ADM_OUTCOME",
            weightCol="classWeights", \
                maxIter=10, regParam=reg_param,
            elasticNetParam=0)
    elif method=="Lasso":
```

```

        lr = LogisticRegression(featuresCol="features", labelCol="ADM_OUTCOME", \
    ↪weightCol="classWeights", \
                                maxIter=10, regParam=reg_param, \
    ↪elasticNetParam=1)

    # Fit the model
    admModel = lr.fit(training_set)

    # predict on testing set
    predict_test=admModel.transform(testing_set)

    # make evaluator
    evaluator=BinaryClassificationEvaluator(rawPredictionCol="rawPrediction", \
    ↪labelCol="ADM_OUTCOME")

    return evaluator.evaluate(predict_test)

```

1.7 Determine best hyperparameters

```

[13]: # test standard LR model
print("ROC-AUC for standard LR is: ", lr_nhamcs(training,testing))

```

ROC-AUC for standard LR is: 0.7540790183387306

```

[14]: # test Ridge LR model for different values of the regularizatoin parameter
for i in np.arange(0.0, 1.1, 0.1):
    i = np.round(i,1)
    print("ROC-AUC for Ridge LR with reg_param=", i, \
          " is: ", lr_nhamcs(training,testing, i,"Ridge"))

```

```

ROC-AUC for Ridge LR with reg_param= 0.0 is: 0.7540790183387306
ROC-AUC for Ridge LR with reg_param= 0.1 is: 0.7540419363538331
ROC-AUC for Ridge LR with reg_param= 0.2 is: 0.7540756472491944
ROC-AUC for Ridge LR with reg_param= 0.3 is: 0.7542239751887837
ROC-AUC for Ridge LR with reg_param= 0.4 is: 0.754362189859765
ROC-AUC for Ridge LR with reg_param= 0.5 is: 0.7542273462783199
ROC-AUC for Ridge LR with reg_param= 0.6 is: 0.75416666666666701
ROC-AUC for Ridge LR with reg_param= 0.7 is: 0.7541026159654836
ROC-AUC for Ridge LR with reg_param= 0.8 is: 0.7540756472491945
ROC-AUC for Ridge LR with reg_param= 0.9 is: 0.7544869201726022
ROC-AUC for Ridge LR with reg_param= 1.0 is: 0.7550094390507027

```

```

[15]: # test Lasso LR model for different values of the regularizatoin parameter
for i in np.arange(0.0, 1.1, 0.1):
    i = np.round(i,1)
    print("ROC-AUC for Lasso LR with reg_param=", i, \

```

```
" is: ", lr_nhamcs(training,testing, i,"Lasso"))
```

```
ROC-AUC for Lasso LR with reg_param= 0.0 is: 0.7540790183387306
ROC-AUC for Lasso LR with reg_param= 0.1 is: 0.7391906014023741
ROC-AUC for Lasso LR with reg_param= 0.2 is: 0.5
ROC-AUC for Lasso LR with reg_param= 0.3 is: 0.5
ROC-AUC for Lasso LR with reg_param= 0.4 is: 0.5
ROC-AUC for Lasso LR with reg_param= 0.5 is: 0.5
ROC-AUC for Lasso LR with reg_param= 0.6 is: 0.5
ROC-AUC for Lasso LR with reg_param= 0.7 is: 0.5
ROC-AUC for Lasso LR with reg_param= 0.8 is: 0.5
ROC-AUC for Lasso LR with reg_param= 0.9 is: 0.5
ROC-AUC for Lasso LR with reg_param= 1.0 is: 0.5
```

1.8 Size of saved model

```
[16]: # create model with best hyperparameters (Ridge, regParam=1.0)
lr = LogisticRegression(featuresCol="features", labelCol="ADM_OUTCOME",
    ↪weightCol="classWeights", \
                           maxIter=10, regParam=1.0, elasticNetParam=0)

admModel = lr.fit(training)
```

```
[17]: # save model
admModel.write().overwrite().save("../models/001-log_regress-no_RFV")
```

```
[18]: # get size on disk
!du -h ../models/001-log_regress-no_RFV
```

```
28K    ../models/001-log_regress-no_RFV/data
20K    ../models/001-log_regress-no_RFV/metadata
52K    ../models/001-log_regress-no_RFV
```

1.9 Get evaluation metrics

```
[19]: # predict on testing set
predict_test=admModel.transform(testing)
```

```
[20]: # calculate AUC
evaluator=BinaryClassificationEvaluator(rawPredictionCol="rawPrediction",
    ↪labelCol="ADM_OUTCOME")
print("ROC-AUC:", evaluator.evaluate(predict_test))
```

```
ROC-AUC: 0.7550094390507027
```

```
[21]: # calculate accuracy
print("F1 score:",evaluator.setMetricName("areaUnderPR").evaluate(predict_test))
```

F1 score: 0.18282251646509748

```
[22]: # calculate accuracy
correct = predict_test.where('prediction == ADM_OUTCOME').count()
total = predict_test.count()

print("Accuracy:", correct/total)
```

Accuracy: 0.7132486388384754

```
[24]: # compute confusion matrix
tp = predict_test.where('prediction == 1 and ADM_OUTCOME==1').count()
fp = predict_test.where('prediction == 1 and ADM_OUTCOME==0').count()
tn = predict_test.where('prediction == 0 and ADM_OUTCOME==0').count()
fn = predict_test.where('prediction == 0 and ADM_OUTCOME==1').count()

print("\nConfusion Matrix:")
print('tn:',tn,' fn:',fn)
print('fp:',fp, ' tp:',tp,)
```

Confusion Matrix:

tn: 1473 fn: 45

fp: 587 tp: 99

```
[ ]: # convert to PDF
!jupyter nbconvert --to pdf `pwd`/DS5559-Final_proj-NHAMCS-Report_2.ipynb
```

```
[ ]:
```